

Loan Eligibility with Ensemble Boosting

Table of Contents

1. Abstract

2. Import of the Dataset

- 2.1. Dataset Description
- 2.2. Missing Values
- 2.3. Numerical and Ordinal Data
- 2.4. General Categorical and Binary Categorical Data

3. Data Preprocessing

4. Training the Models

- 4.1. eXtreme Gradient Boosting
- 4.2. Adaptive Boosting
- 4.3. Gradient Boosting

5. Comparisons

6. Conclusions and Possible Expansions

7. Appendix

- 7.A. Python stay-awake module
- 7.B. Python image converter base64 (PDF exporting)

1. Abstract

Banks and all sort of credit institutes which provide *loans*, face great challenges when a new potential customer asks for a loan.

These challenges translate into *risks*, which if not properly balanced, may end up with catastrophic consequences for all clients in the financial-chain, and may have a big impact in the outside world too.

Technology can mitigate these risks.

By gathering a huge amount of *data*, and by applying some analytics we can see the *hidden shape* of the risks that credit institutes take.

Then, thanks to some state-of-the-art algorithms we can predict the potential risk that a new client may represent for a bank.

At the end of the day, the aim is to provide a robust model to support difficult *decision-making* scenarios.

The purpose of the notebook is to show and manipulate a huge dataset collected from [Kaggle](#), then after some *feature engineering* we apply and compare three State-of-the-Art *ensemble learning* algorithms for tabular data:

- AdaBoost
- GradientBoost
- XGBoost

Based on statistics we decide which the best one is.

Further conclusions and possible expansions are presented.

[pointer to the dataset on Kaggle](#)



See [Appendix 7.B](#). to understand how this image has been generated

2. Dataset Import

```
In [1]: #import of all libraries and packages
import kagglehub #for automatically downloading the dataset from Kaggle IF it has c
import pandas as pd #data manipulation
import numpy as np #numerical manipulation
import matplotlib.pyplot as plt #data visualization
import matplotlib.patches as mpatches #for distinguishing the dots in the scatter m
import sklearn #ML
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder #handling categoric
import sys #get the real size of the DataFrames
```

```
In [2]: #Code to fetch the data
try:
    #Download latest version
    path = kagglehub.dataset_download("yasserh/loan-default-dataset")

    #to check where the dataset has been located
```

```
#print("Path to dataset files:", path)

loan_data = pd.read_csv(path + '\Loan_Default.csv')

#if something goes wrong (e.g.: connection, wrong path, ...), use the copy of the d
except:
    loan_data = pd.read_csv('Loan_Default.csv')
```

Warning: Looks like you're using an outdated `kagglehub` version (installed: 0.3.12), please consider upgrading to the latest version (0.3.13).

2.1. Dataset Description

```
In [3]: pd.set_option('display.max_columns',
                    None) #display all columns
loan_data.head()
```

```
Out[3]:
```

	ID	year	loan_limit	Gender	approv_in_adv	loan_type	loan_purpose	Credit_Worth
0	24890	2019	cf	Sex Not Available	nopre	type1	p1	
1	24891	2019	cf	Male	nopre	type2	p1	
2	24892	2019	cf	Male	pre	type1	p1	
3	24893	2019	cf	Male	nopre	type1	p4	
4	24894	2019	cf	Joint	pre	type1	p1	

This description may be very helpful through all the notebook's reading, so keep it in sight.

- **ID:** client loan application id
- **year:** year of loan application
- **loan limit:** indicates whether the loan is conforming (cf) or non-conforming (ncf)
- **Gender:** gender of the applicant (male, female, joint, sex not available)
- **approv_in_adv:** indicates whether the loan was approved in advance (pre, nopre)
- **loan_type:** type of loan (type1, type2, type3) :
 - *Type 1 (Conventional Loans):* Characterized by higher loan amounts, lower LTV ratios, and stronger credit scores, making them a preferred option for well-qualified, lower-risk borrowers.
 - *Type 2 (Government-Backed Loans):* Typically involve lower loan amounts, higher LTV ratios, and moderate credit scores, indicating they are used by borrowers with smaller down payments who benefit from government-backed programs.
 - *Type 3 (Non-Conventional Loans):* Feature moderate loan amounts, the highest LTV ratios, and lower credit scores, often associated with higher-risk products such as

jumbo loans or adjustable-rate mortgages.

- **loan_purpose:** purpose of the loan (p1, p2, p3, p4):
 - *p1 (Home Purchase):* Represents loans taken out for primary residences, often displaying moderate credit scores and higher LTV ratios.
 - *p2 (Home Improvement):* Smaller loan amounts used for property renovations, with lower LTV ratios suggesting homeowners are leveraging built-up equity.
 - *p3 (Refinancing):* Applies to homeowners replacing an existing mortgage, characterized by moderate loan amounts and lower LTV ratios, indicating financial stability.
 - *p4 (Investment Property):* Involves larger loan amounts and higher risk profiles, primarily financed through conventional loans due to restrictions on Government-backed funding for investment properties.
- **Credit_Worthiness:** credit worthiness (l1, l2)
- **open_credit:** indicates whether the applicant has any open credit accounts (opc, nopc)
- **business_or_commercial:** indicates whether the loan is for business/commercial purposes (ob/c - business/commercial, nob/c - personal)
- **loan_amount:** amount of money being borrowed
- **rate_of_interest:** interest rate charged on the loan
- **Interest_rate_spread:** difference between the interest rate on the loan and a benchmark interest rate
- **Upfront_charges:** initial charges associated with securing the loan
- **term:** duration of the loan in months
- **Neg_ammortization:** indicates whether the loan allows for negative amortization (neg_amm, not_neg)
- **interest_only:** indicates whether the loan has an interest-only payment option (int_only, not_int)
- **lump_sum_payment:** indicates if a lump sum payment is required at the end of the loan term (lpsm, not_lpsm)
- **property_value:** value of the property being financed
- **construction_type:** type of construction (sb - site built, mh - manufactured home)
- **occupancy_type:** occupancy type (pr - primary residence, sr - secondary residence, ir - investment property)
- **Secured_by:** specifies the type of collateral securing the loan (home, land)
- **total_units:** number of units in the property being financed (1U, 2U, 3U, 4U)
- **income *** : applicant's annual income
- **credit_type:** applicant's type of credit (CIB - credit information bureau, CRIF - CIRF credit information bureau, EXP - experian, EQUI - equifax)
- **Credit_Score:** applicant's credit score
- **co-applicant_credit_type:** co-applicant's type of credit (CIB - credit information bureau, EXP - experian)
- **age:** the age of the applicant

- **submission_of_application:** indicates how the application was submitted (to_inst - to institution, not_inst - not to institution)
- **LTV:** loan-to-value ratio, calculated as the loan amount divided by the property value
- **Region:** geographic region where the property is located (North, South, Central, North-East)
- **Security_Type:** type of security or collateral backing the loan (direct, indirect)
- **Status:** indicates whether the loan has been defaulted (1) or not (0)
- **dtir1:** debt-to-income ratio

* The annual income is in thousands of dollars.

```
In [4]: loan_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148670 entries, 0 to 148669
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     148670 non-null  int64
1   year                                 148670 non-null  int64
2   loan_limit                           145326 non-null  object
3   Gender                               148670 non-null  object
4   approv_in_adv                        147762 non-null  object
5   loan_type                            148670 non-null  object
6   loan_purpose                           148536 non-null  object
7   Credit_Worthiness                    148670 non-null  object
8   open_credit                          148670 non-null  object
9   business_or_commercial               148670 non-null  object
10  loan_amount                          148670 non-null  int64
11  rate_of_interest                     112231 non-null  float64
12  Interest_rate_spread                 112031 non-null  float64
13  Upfront_charges                      109028 non-null  float64
14  term                                 148629 non-null  float64
15  Neg_ammortization                    148549 non-null  object
16  interest_only                        148670 non-null  object
17  lump_sum_payment                     148670 non-null  object
18  property_value                       133572 non-null  float64
19  construction_type                    148670 non-null  object
20  occupancy_type                       148670 non-null  object
21  Secured_by                           148670 non-null  object
22  total_units                           148670 non-null  object
23  income                               139520 non-null  float64
24  credit_type                           148670 non-null  object
25  Credit_Score                          148670 non-null  int64
26  co-applicant_credit_type              148670 non-null  object
27  age                                   148470 non-null  object
28  submission_of_application             148470 non-null  object
29  LTV                                   133572 non-null  float64
30  Region                               148670 non-null  object
31  Security_Type                         148670 non-null  object
32  Status                               148670 non-null  int64
33  dtir1                                124549 non-null  float64
dtypes: float64(8), int64(5), object(21)
memory usage: 38.6+ MB

```

Notice: above the `info()` command, gives us a memory usage size that is actually a bit misleading, [GitHub issue DataFrame.memory_usage](#).

That is the reason why we decided to use: `sys.getsizeof()` , since this data will be quite interesting especially for the next sections.

```
In [5]: sys.getsizeof(loan_data) #returns the size in bytes
```

```
Out[5]: 207281850
```

```
In [6]: loan_data.describe()
```

	ID	year	loan_amount	rate_of_interest	Interest_rate_spread	Upfron
count	148670.000000	148670.0	1.486700e+05	112231.000000	112031.000000	1090
mean	99224.500000	2019.0	3.311177e+05	4.045476	0.441656	32
std	42917.476598	0.0	1.839093e+05	0.561391	0.513043	32
min	24890.000000	2019.0	1.650000e+04	0.000000	-3.638000	
25%	62057.250000	2019.0	1.965000e+05	3.625000	0.076000	5
50%	99224.500000	2019.0	2.965000e+05	3.990000	0.390400	25
75%	136391.750000	2019.0	4.365000e+05	4.375000	0.775400	48
max	173559.000000	2019.0	3.576500e+06	8.000000	3.357000	600

Since the columns: **ID** and **year** are not meaningful for any purpose, we can drop them from the very beginning and still be compliant to the best preprocessing practices.

```
In [7]: #since they are not meaningful columns at all, we can drop them from the whole data
loan_data = loan_data.drop('ID', axis=1)
loan_data = loan_data.drop('year', axis=1)
```

```
In [8]: #to get insights on the categorical data
categorical_values_counts = list()
for cat in loan_data.select_dtypes(include=['object']).columns: #extract the list of
    categorical_values_counts.append(loan_data[f"{cat}"].value_counts())
    categorical_values_counts.append('-'*45) #for spacing them
categorical_values_counts
```

```

Out[8]: [cf      135348
ncf      9978
Name: loan_limit, dtype: int64,
'-----',
Male      42346
Joint     41399
Sex Not Available  37659
Female     27266
Name: Gender, dtype: int64,
'-----',
nopre    124621
pre      23141
Name: approv_in_adv, dtype: int64,
'-----',
type1    113173
type2    20762
type3    14735
Name: loan_type, dtype: int64,
'-----',
p3      55934
p4      54799
p1      34529
p2      3274
Name: loan_purpose, dtype: int64,
'-----',
l1      142344
l2       6326
Name: Credit_Worthiness, dtype: int64,
'-----',
nopc    148114
opc       556
Name: open_credit, dtype: int64,
'-----',
nob/c    127908
b/c      20762
Name: business_or_commercial, dtype: int64,
'-----',
not_neg   133420
neg_amm    15129
Name: Neg_ammortization, dtype: int64,
'-----',
not_int   141560
int_only    7110
Name: interest_only, dtype: int64,
'-----',
not_lpsm   145286
lpsm       3384
Name: lump_sum_payment, dtype: int64,
'-----',
sb      148637
mh        33
Name: construction_type, dtype: int64,
'-----',
pr      138201
ir       7340
sr       3129

```



```

Name: occupancy_type, dtype: int64,
'-----',
home      148637
land       33
Name: Secured_by, dtype: int64,
'-----',
1U      146480
2U       1477
3U       393
4U       320
Name: total_units, dtype: int64,
'-----',
CIB      48152
CRIF     43901
EXP      41319
EQUI     15298
Name: credit_type, dtype: int64,
'-----',
CIB      74392
EXP      74278
Name: co-applicant_credit_type, dtype: int64,
'-----',
45-54     34720
35-44     32818
55-64     32534
65-74     20744
25-34     19142
>74       7175
<25       1337
Name: age, dtype: int64,
'-----',
to_inst    95814
not_inst   52656
Name: submission_of_application, dtype: int64,
'-----',
North      74722
south      64016
central     8697
North-East  1235
Name: Region, dtype: int64,
'-----',
direct     148637
Indriect    33
Name: Security_Type, dtype: int64,
'-----']

```

2.2. Missing Values

We can see that the biggest issues in terms of missing values come from:

Upfront_charges , **Interest_rate_spread** , **rate_of_interest** . we could have some criticalities also on other features but still negligible, even though we have to keep in mind that the **NaN** are spread out as:

Numerical

- 39642 Upfront_charges
- 36639 Interest_rate_spread
- 36439 rate_of_interest
- 24121 dtir1
- 15098 LTV
- 15098 property_value
- 9150 income
- 41 term

Categorical

- 3344 loan_limit
- 908 approv_in_adv
- 200 age
- 200 submission_of_application
- 136 loan_purpose
- 121 Neg_ammortization

In the next code blocks we will see how to approach this problem.

2.3. Numerical and Ordinal Data

```
In [9]: #converting 'total_units' and 'age' categorical data into ordinal
ordinal_encoder = OrdinalEncoder()

#just for the total_units and age features, since it makes sense a concept of ordinal
total_units_cat = loan_data[['total_units']] #take the column we want to encode
total_units_encoded = ordinal_encoder.fit_transform(total_units_cat) + 1 #we shift
loan_data['total_units'] = total_units_encoded #replace the column in the dataset

age_mapping = {
    '<25':1,
    '25-34':2,
    '35-44':3,
    '45-54':4,
    '55-64':5,
    '65-74':6,
    '>74':7
}
loan_data['age'] = loan_data['age'].map(age_mapping) #replace the column in the dataset
```


```
In [10]: #list containing all the numerical + ordinal features plus the target
numerical_ordinal_columns = [col for col in loan_data.columns if col not in loan_data.columns]
#list with only the numerical + ordinal features
numerical_ordinal_features = numerical_ordinal_columns.copy()
numerical_ordinal_features.remove('Status')
```

```
In [11]: #boxplots
loan_data[numerical_ordinal_features].plot(kind='box',
```

```

subplots=True, #distinct box plots for each feature
layout= (int(len(numerical_ordinal_features) ** 0.5)
figsize=((int(len(numerical_ordinal_features) ** 0.5
sharex=False, #each subplot uses a different x scale
sharey=False #each subplot uses a different y scale
)
plt.show()

```



The box plots above show that in several occasions we encounter outliers.

Remember from statistics that:

$$\text{Outlier Range} = [\text{Q}_1 - 1.5 \cdot \text{IQR}, \text{Q}_3 + 1.5 \cdot \text{IQR}]$$


Where:

- Q_1 : first quartile
- Q_3 : third quartile
- InterQuartile Range: $\text{IQR} = \text{Q}_3 - \text{Q}_1$

```

In [12]: #plot the histograms
loan_data.hist(bins=50, #number of bins to encapsulate the data
               figsize=(20,15)
               );

```



Notice: There are some capped values, specifically concerning:

- term
- dtir1
- Credit_Score

term feature for sure is capped at: \$360\$. Probably also: **dtir1** and **Credit_Score**, were capped but with a lower impact.

In any case, since they are not our target attributes, this is not a big issue and we can actually ignore it.

```
In [13]: loan_data.corr()["Status"].sort_values(ascending=False) #correlation related to out
```

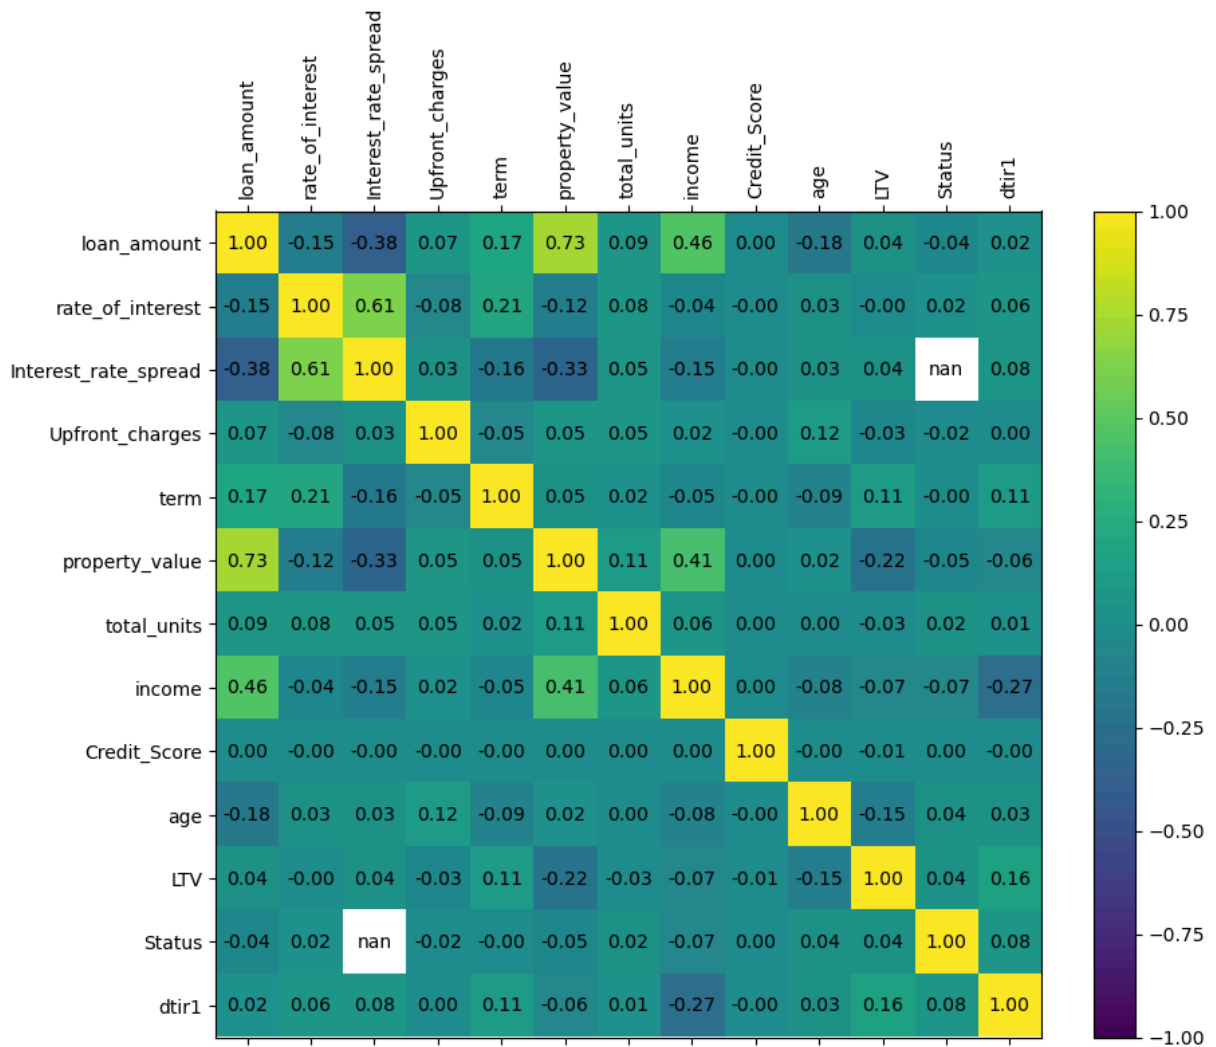
```
Out[13]: Status          1.000000
dtir1                0.078083
age                  0.044600
LTV                  0.038895
total_units          0.023800
rate_of_interest     0.022957
Credit_Score         0.004004
term                 -0.000240
Upfront_charges      -0.019138
loan_amount          -0.036825
property_value        -0.048864
income               -0.065119
Interest_rate_spread      NaN
Name: Status, dtype: float64
```

```
In [14]: #compute the correlations
correlations = loan_data.corr()

#plot
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111)
cax = ax.matshow(correlations,
                  vmin=-1,
                  vmax=1)
fig.colorbar(cax)
ticks = np.arange(len(numerical_ordinal_columns))
ax.set_xticks(ticks)
ax.set_yticks(ticks)
ax.set_xticklabels(numerical_ordinal_columns,
                   rotation=90) #to see the column labels clearly
ax.set_yticklabels(numerical_ordinal_columns)

# Annotate cells with correlation values
for i in range(len(correlations)):
    for j in range(len(correlations)):
        value = correlations.iloc[i, j]
        ax.text(j,
                i,
                f'{value:.2f}',
                va='center',
                ha='center',
                color='black')

plt.show()
```



Notice: the correlation coefficient between **Status** and **interest_rate_spread** returns not a number, this is due to the fact that the latter feature is actually compromised, (we will see in few code blocks how), furthermore the two standard deviations are very tiny.

```
In [15]: #showing the percentages of data grouped by target values
def compute_feature_target_percentages(df, features, target='Status'):
    rows = []
    for feat in features:
        #Discharge NaN values for feature and target
        feature_data = df[[feat, target]].dropna()

        #total valid values for this feature
        total = len(feature_data)

        #How many of these are target=0 and target=1
        count_0 = feature_data[feature_data[target] == 0].shape[0]
        count_1 = feature_data[feature_data[target] == 1].shape[0]

        #we translate the above data into percentages
        pct_0 = (count_0 / total) * 100 if total > 0 else 0
        pct_1 = (count_1 / total) * 100 if total > 0 else 0
```

```

        rows.append({'feature': feat,
                     'target=0 %': pct_0,
                     'target=1 %': pct_1})
    return pd.DataFrame(rows)

result = compute_feature_target_percentages(loan_data, numerical_ordinal_features)

#Extract data and labels
data_matrix = result[['target=0 %',
                      'target=1 %']].to_numpy()
feature_labels = result['feature'].tolist()
target_labels = ['Status 0',
                 'Status 1']


#Plot
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111)
cax = ax.matshow(data_matrix,
                 vmin=0,
                 vmax=100,
                 cmap='coolwarm_r')
fig.colorbar(cax)

#Ticks
ax.set_xticks(np.arange(len(target_labels)))
ax.set_yticks(np.arange(len(feature_labels)))
ax.set_xticklabels(target_labels,
                  rotation=90)
ax.set_yticklabels(feature_labels)

#Values
for i in range(data_matrix.shape[0]):
    for j in range(data_matrix.shape[1]):
        value = data_matrix[i, j]
        ax.text(j,
                i, f'{value:.1f}%',
                va='center',
                ha='center',
                color='black')

plt.show()

```



Notice: it is important to point out that through the notebook we adopted the following convention:

- **Status=0** , i.e.: loan totally repaid, is chosen as the **negative** case
- **Status=1** , i.e.: loan not repaid, is chosen as the **positive** case

This is due to the general rule that the *positive* class should be the one of interest, so the event that we want to detect, indeed for us it is important to detect the default client.

Due to what we saw in section: [2.2. Missing Values](#), we can say that: **rate_of_interest** , **Interest_rate_spread** and **Upfront_charges** , when they are not missing they are

nearly always with **Status=0** . We must drop them entirely from the whole dataset, because they are not meaningful and they have also the highest misses.

```
In [16]: #discharge the not desired columns
data = loan_data.drop(['rate_of_interest',
                      'Interest_rate_spread',
                      'Upfront_charges'],
                      axis=1)


#discharge the rows without elements
data = data.dropna()

#prepare the colors based on the Status (0 o 1)
color_map = {0: 'blue',
             1: 'red'
            }
colors = data['Status'].map(color_map)

# Plot of the scatter_matrix
pd.plotting.scatter_matrix(data.drop(columns=['Status']),
                           figsize=(20, 15),
                           diagonal='hist',
                           color=colors,
                           alpha=0.5
                           )

#patch for the Legend
legend_handles = [mpatches.Patch(color='blue',
                                  label='Status = 0'
                                  ),
                  mpatches.Patch(color='red',
                                  label='Status = 1'
                                  )
                  ]

plt.legend(handles=legend_handles,
           loc='upper right',
           bbox_to_anchor=(1.15, -0.3)
           )
plt.show()
```



Notice: grouping dots based on the the target condition is very helpful for data visualization, there is a Python library for data visualization [Seaborn](#) * that allows to plot directly these types of graphs in a very straightforward way, but it is a more computational expensive tool, that in cases of big datasets can slow down the process a lot, this is the reason why we implemented the solution using `matplotlib` in a longer way.

In some cells of the matrix, the color division makes us able to draw a straight line that can divide the default and non-default loans in an approximately sharp manner.

*: [Here](#) is a project of mine with an implention using Seaborn

General Categorical and Binary Categorical Data

Notice: in the following block of code, we decided to transform manually the binary categorical data in binary ordinal data, assigning \$0\$ or \$1\$ with the following concept: \$0\$ whether the associated variable has lower probability of deaault, \$1\$ viceversa. (where possible).

```
In [17]: #convert all the categorical binary columns to ordinal binary
loan_data['loan_limit'] = loan_data['loan_limit'].map({'cf': 1, 'ncf': 0})
loan_data['approv_in_adv'] = loan_data['approv_in_adv'].map({'nopre': 1, 'pre': 0})
loan_data['Credit_Worthiness'] = loan_data['Credit_Worthiness'].map({'l2': 1, 'l1': 0})
loan_data['open_credit'] = loan_data['open_credit'].map({'opc': 1, 'nopc': 0})
```

```

loan_data['business_or_commercial'] = loan_data['business_or_commercial'].map({'b/c': 1, 'not_b/c': 0})
loan_data['Neg_ammortization'] = loan_data['Neg_ammortization'].map({'not_neg': 1, 'neg': 0})
loan_data['interest_only'] = loan_data['interest_only'].map({'not_int': 1, 'int_only': 0})
loan_data['lump_sum_payment'] = loan_data['lump_sum_payment'].map({'not_lpsm': 1, 'lpsm': 0})
loan_data['construction_type'] = loan_data['construction_type'].map({'sb': 1, 'mh': 0})
loan_data['Secured_by'] = loan_data['Secured_by'].map({'home': 1, 'land': 0})
loan_data['co-applicant_credit_type'] = loan_data['co-applicant_credit_type'].map({'co_applicant': 1, 'not_co_applicant': 0})
loan_data['submission_of_application'] = loan_data['submission_of_application'].map({'submission': 1, 'no_submission': 0})
loan_data['Security_Type'] = loan_data['Security_Type'].map({'direct': 1, 'Indirect': 0})

```

```

In [18]: #List of categorical non-binary columns
cat_columns = ['Gender',
               'loan_type',
               'loan_purpose',
               'occupancy_type',
               'credit_type',
               'Region']

# Loop through each column
for col in cat_columns:
    encoder = OneHotEncoder(sparse=False,
                           handle_unknown='ignore') #new encoder for each
    encoded_array = encoder.fit_transform(loan_data[[col]])
    encoded_df = pd.DataFrame(encoded_array,
                              columns=encoder.get_feature_names_out([col]),
                              index=loan_data.index)

# Drop original column and concatenate the encoded DataFrame
loan_data.drop(columns=[col],
                inplace=True)
loan_data = pd.concat([loan_data,
                       encoded_df
                       ],
                      axis=1)

```

```

In [19]: #List of all categorical columns
categorical_columns = [
    'loan_limit',
    'approv_in_adv',
    'Credit_Worthiness',
    'open_credit',
    'business_or_commercial',
    'Neg_ammortization',
    'interest_only',
    'lump_sum_payment',
    'construction_type',
    'Secured_by',
    'co-applicant_credit_type',
    'submission_of_application',
    'Security_Type',
    'Gender_Male', 'Gender_Joint', 'Gender_Sex Not Available', 'Gender_Sex Not Available',
    'loan_type_type1', 'loan_type_type2', 'loan_type_type3',
    'loan_purpose_p3', 'loan_purpose_p4', 'loan_purpose_p1', 'loan_purpose_p2',
    'occupancy_type_pr', 'occupancy_type_ir', 'occupancy_type_sr',
    'credit_type_CIB', 'credit_type_CRIF', 'credit_type_EXP', 'credit_type_IR',

```

```

        'Region_North', 'Region_south', 'Region_central', 'Region_Nor
        'Status'
    ]

categorical_features = categorical_columns.copy()
#list of all categorical features
categorical_features.remove('Status')

```

```

In [20]: binary_correlation = loan_data[categorical_columns].corr()


#plot
fig = plt.figure(figsize=(20, 16))
ax = fig.add_subplot(111)
cax = ax.matshow(binary_correlation,
                  vmin=-1,
                  vmax=1)

ticks = np.arange(len(categorical_columns))
ax.set_xticks(ticks)
ax.set_yticks(ticks)
ax.set_xticklabels(categorical_columns,
                  rotation=90) #to see the column lables clearly
ax.set_yticklabels(categorical_columns)

#annotate cells with correlation values
for i in range(len(binary_correlation)):
    for j in range(len(binary_correlation)):
        value = binary_correlation.iloc[i, j]
        ax.text(j,
                i,
                f'{value:.2f}',
                va='center',
                ha='center',
                color='black')

plt.show()

```



Pandas `dataframe.corr()` method computes the Pearson's correlation coefficient:

$$\rho_{x,y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y}$$

where:

- cov is the covariance computed as: $\text{cov}(X,Y) = E[(X - \mu_X)(Y - \mu_Y)]$
- σ_X is the standard deviation of X
- σ_Y is the standard deviation of Y

In the case of binary variables that assume only values in $\{0,1\}$, using Pearson's coefficient is a possible way to gain an overview of the correlation between variables

[ResearchGate correlation of binary data.](#)

Keep in mind that in literature there exist many other methods to achieve so, such as: point-biserial correlation coefficient, Spearman's correlation coefficient.

```

In [21]: def compute_binary_feature_target_distribution(df, binary_features, target='Status')
    rows = []
    for feat in binary_features:
        for val in [0, 1]:
            #filter for feature value
            subset = df[(df[feat] == val) & (df[target].isin([0, 1]))]

            total = len(subset)
            count_0 = subset[subset[target] == 0].shape[0]
            count_1 = subset[subset[target] == 1].shape[0]

            pct_0 = (count_0 / total) * 100 if total > 0 else 0
            pct_1 = (count_1 / total) * 100 if total > 0 else 0

            rows.append({
                'feature': feat,
                'value': val,
                'Status 0 %': pct_0,
                'Status 1 %': pct_1
            })
    return pd.DataFrame(rows)

result_bin = compute_binary_feature_target_distribution(loan_data, categorical_feat

#extract data for heatmap
data_matrix = result_bin[['Status 0 %', 'Status 1 %']].to_numpy()
row_labels = [f"{f} = {v}" for f, v in zip(result_bin['feature'], result_bin['value'])]
col_labels = ['Status 0', 'Status 1']

#plot
fig = plt.figure(figsize=(15, len(row_labels) * 0.7))
ax = fig.add_subplot(111)
cax = ax.matshow(data_matrix,
                  vmin=0,
                  vmax=100,
                  cmap='coolwarm_r')

# ticks
ax.set_xticks(np.arange(len(col_labels)))
ax.set_yticks(np.arange(len(row_labels)))
ax.set_xticklabels(col_labels, rotation=90)
ax.set_yticklabels(row_labels)

#annotate values
for i in range(data_matrix.shape[0]):
    for j in range(data_matrix.shape[1]):
        value = data_matrix[i, j]
        ax.text(j,
                i,
                f'{value:.1f}%',
                va='center',
                ha='center',
                color='black')

plt.show()

```

	Status 0	Status 1
loan_limit = 0	66.8%	33.2%
loan_limit = 1	76.0%	24.0%
approv_in_adv = 0	79.1%	20.9%
approv_in_adv = 1	74.7%	25.3%
Credit_Worthiness = 0	75.7%	24.3%
Credit_Worthiness = 1	68.2%	31.8%
open_credit = 0	75.3%	24.7%
open_credit = 1	82.4%	17.6%
business_or_commercial = 0	77.0%	23.0%
business_or_commercial = 1	65.5%	34.5%
Neg_ammortization = 0	55.4%	44.6%
Neg_ammortization = 1	77.6%	22.4%
interest_only = 0	72.7%	27.3%
interest_only = 1	75.5%	24.5%
lump_sum_payment = 0	22.3%	77.7%
lump_sum_payment = 1	76.6%	23.4%

construction_type = 0	0.0%	100.0%
construction_type = 1	75.4%	24.6%
Secured_by = 0	0.0%	100.0%
Secured_by = 1	75.4%	24.6%
co-applicant_credit_type = 0	69.1%	30.9%
co-applicant_credit_type = 1	81.6%	18.4%
submission_of_application = 0	82.5%	17.5%
submission_of_application = 1	71.6%	28.4%
Security_Type = 0	0.0%	100.0%
Security_Type = 1	75.4%	24.6%
Gender_Male = 0	76.0%	24.0%
Gender_Male = 1	73.8%	26.2%
Gender_Joint = 0	73.2%	26.8%
Gender_Joint = 1	80.8%	19.2%
Gender_Sex Not Available = 0	76.7%	23.3%
Gender_Sex Not Available = 1	71.4%	28.6%
Gender_Female = 0	75.5%	24.5%
Gender_Female = 1	74.9%	25.1%

loan_type_type1 = 0	69.4%	30.6%
loan_type_type1 = 1	77.2%	22.8%
loan_type_type2 = 0	77.0%	23.0%
loan_type_type2 = 1	65.5%	34.5%
loan_type_type3 = 0	75.4%	24.6%
loan_type_type3 = 1	74.9%	25.1%
loan_purpose_p3 = 0	75.6%	24.4%
loan_purpose_p3 = 1	75.0%	25.0%
loan_purpose_p4 = 0	74.4%	25.6%
loan_purpose_p4 = 1	77.0%	23.0%
loan_purpose_p1 = 0	75.7%	24.3%
loan_purpose_p1 = 1	74.1%	25.9%
loan_purpose_p2 = 0	75.5%	24.5%
loan_purpose_p2 = 1	66.9%	33.1%
occupancy_type_pr = 0	70.9%	29.1%
occupancy_type_pr = 1	75.7%	24.3%
occupancy_type_ir = 0	75.6%	24.4%
occupancy_type_ir = 1	70.0%	30.0%

occupancy_type_sr = 0	75.4%	24.6%
occupancy_type_sr = 1	72.9%	27.1%
credit_type_CIB = 0	71.1%	28.9%
credit_type_CIB = 1	84.2%	15.8%
credit_type_CRIF = 0	71.8%	28.2%
credit_type_CRIF = 1	83.8%	16.2%
credit_type_EXP = 0	72.0%	28.0%
credit_type_EXP = 1	84.0%	16.0%
credit_type_EQUI = 0	84.0%	16.0%
credit_type_EQUI = 1	0.0%	100.0%
Region_North = 0	73.2%	26.8%
Region_North = 1	77.5%	22.5%
Region_south = 0	76.9%	23.1%
Region_south = 1	73.4%	26.6%
Region_central = 0	75.5%	24.5%
Region_central = 1	72.5%	27.5%
Region_North-East = 0	75.4%	24.6%
Region_North-East = 1	69.6%	30.4%



3. Data Preprocessing

So far we have handled the whole dataset manually in order to better manipulate it and visualize relevant characteristics.

In machine learning, it is way better to use **pipelines**, which are automated and organized series of steps that orchestrate the entire machine learning lifecycle, from data collection and preprocessing to model training, validation, and deployment.

```
In [22]: #import of packages and libraries useful for fitting, transforming hence pipelines
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import FunctionTransformer
from sklearn.impute import SimpleImputer
import warnings #for ignoring the warnings while the models run

In [23]: #while running some code blocks we encounter some useless warnings that in producti
#ignore the FutureWarning
warnings.filterwarnings("ignore", category=FutureWarning)
#ignore the UserWarning
warnings.filterwarnings("ignore", category=UserWarning)

In [24]: #features are now only numerical and ordinal, the number of them increased, as the
loan_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148670 entries, 0 to 148669
Data columns (total 49 columns):
```

#	Column	Non-Null	Count	Dtype
0	loan_limit	145326	non-null	float64
1	approv_in_adv	147762	non-null	float64
2	Credit_Worthiness	148670	non-null	int64
3	open_credit	148670	non-null	int64
4	business_or_commercial	148670	non-null	int64
5	loan_amount	148670	non-null	int64
6	rate_of_interest	112231	non-null	float64
7	Interest_rate_spread	112031	non-null	float64
8	Upfront_charges	109028	non-null	float64
9	term	148629	non-null	float64
10	Neg_ammortization	148549	non-null	float64
11	interest_only	148670	non-null	int64
12	lump_sum_payment	148670	non-null	int64
13	property_value	133572	non-null	float64
14	construction_type	148670	non-null	int64
15	Secured_by	148670	non-null	int64
16	total_units	148670	non-null	float64
17	income	139520	non-null	float64
18	Credit_Score	148670	non-null	int64
19	co-applicant_credit_type	148670	non-null	int64
20	age	148470	non-null	float64
21	submission_of_application	148470	non-null	float64
22	LTV	133572	non-null	float64
23	Security_Type	148670	non-null	int64
24	Status	148670	non-null	int64
25	dtir1	124549	non-null	float64
26	Gender_Female	148670	non-null	float64
27	Gender_Joint	148670	non-null	float64
28	Gender_Male	148670	non-null	float64
29	Gender_Sex Not Available	148670	non-null	float64
30	loan_type_type1	148670	non-null	float64
31	loan_type_type2	148670	non-null	float64
32	loan_type_type3	148670	non-null	float64
33	loan_purpose_p1	148670	non-null	float64
34	loan_purpose_p2	148670	non-null	float64
35	loan_purpose_p3	148670	non-null	float64
36	loan_purpose_p4	148670	non-null	float64
37	loan_purpose_nan	148670	non-null	float64
38	occupancy_type_ir	148670	non-null	float64
39	occupancy_type_pr	148670	non-null	float64
40	occupancy_type_sr	148670	non-null	float64
41	credit_type_CIB	148670	non-null	float64
42	credit_type_CRIF	148670	non-null	float64
43	credit_type_EQUI	148670	non-null	float64
44	credit_type_EXP	148670	non-null	float64
45	Region_North	148670	non-null	float64
46	Region_North-East	148670	non-null	float64
47	Region_central	148670	non-null	float64
48	Region_south	148670	non-null	float64

```
dtypes: float64(37), int64(12)
memory usage: 55.6 MB
```

Notice: the memory usage given directly by Pandas has increased by 17+ MB, due to the transformation into ordinal data.

In order to have a clearer idea of the actual memory used, as already mentioned in section [2.1. Dataset Description](#), we should use `sys.getsizeof()` as below.

```
In [25]: #the result is in bytes  
sys.getsizeof(loan_data)
```

```
Out[25]: 58278784
```

The number shows that actually the memory used is way less with respect to the beginning, this is due to the fact that integer objects take less memory compared to string objects.

Since we want to embrace the *pipeline* approach, **we drop the previously manually manipulated dataset, and we take a new one, as it follows.**

```
In [26]: try:  
         loan_data = pd.read_csv(path + '\Loan_Default.csv')  
  
     except:  
         loan_data = pd.read_csv('Loan_Default.csv') #if something goes wrong. use the c
```

In the following code block we apply column removal, based on the characteristics that emerged in sections: [2.3.](#) and [2.4.](#).

It is important to point out that the column removal has been applied on the whole dataset, because the criticalities on those features were so strong that allowed us to proceed without compromising the effectiveness of the test set.

In order to better explain how the decisions have been taken we divide the data into numerical and categorical:

- **Numerical:**

- `ID` , `year` : have been removed since they did not bring any relevant information for predictive purposes.
- `rate_of_interest` , `Interest_rate_spread` , `Upfront_charges` : have been removed, because they were strongly compromised, they have the highest missing values and furthermore whenever they are not missing, the target is nearly always: `Status=0`

- **Categorical:**

- `business_or_commercial` : has been removed because from the correlation matrix emerges that is strongly correlated to `loan_type_2`
- `Secured_by` , `Security_Type` : have been removed because from the correlation matrix emerged that they were strongly correlated to each other and to `construction_type` , furthermore whenever `Secured_by=land` or `Security_Type=indirect` the target always assumes: `Status=1`

- When the feature `credit_type` assumes the value `EQUI`, the target is always `Status=1`, in this situation it is spuriously predictive. It is a case of [data leakage](#), by removing it, we prevent the model from learning this allegedly wrong pattern. `credit_type` is later transformed with one-hot encoding, because it can assume four different values, so we decided to set it to `NaN` whenever it assumes `EQUI`. Later, when the models call the *fit* of the *pipeline* on only the training set, the `NaN` will be replaced with the most frequent category.

It is important to point out that the threshold of correlation by which we decided to remove columns has been set to 0.90, since decision trees and even better ensemble learning algorithms are quite robust when they face multicollinearity, as opposed to other simpler models such as logistic regression.

```
In [27]: #removing the non-relevant features
loan_data = loan_data.drop([#numerical
                             'ID',
                             'year',
                             'rate_of_interest',
                             'Interest_rate_spread',
                             'Upfront_charges',
                             #categorical
                             'Security_Type',
                             'Secured_by',
                             'business_or_commercial',
                             ],
                             axis=1
                             )

#setting to NaN the non-relevant columns
loan_data.loc[loan_data['credit_type'] == 'EQUI', 'credit_type'] = np.nan
```

```
In [28]: #map for the ordinal mapping of age groups
age_mapping = {'<25': 1,
               '25-34': 2,
               '35-44': 3,
               '45-54': 4,
               '55-64': 5,
               '65-74': 6,
               '>74': 7
               }

#map for the ordinal mapping of the total_units
total_units_mapping = {'1U': 1,
                       '2U': 2,
                       '3U': 3,
                       '4U': 4
                       }

#List of the names of the ordinal features
ordinal_features = ['age',
                    'total_units',
                    ]

#List of the names of the categorical features where we will apply OneHotEncoding
```

```

onehot_features = ['loan_limit',
                  'approv_in_adv',
                  'Credit_Worthiness',
                  'open_credit',
                  'Neg_ammortization',
                  'interest_only',
                  'lump_sum_payment',
                  'construction_type',
                  'co-applicant_credit_type',
                  'submission_of_application',
                  'Gender',
                  'loan_type',
                  'loan_purpose',
                  'occupancy_type',
                  'credit_type',
                  'Region',
                  ]

#list of the numerical features
numerical_features = ['loan_amount',
                     'term',
                     'property_value',
                     'income',
                     'Credit_Score',
                     'LTV',
                     'dtir1'
                     ]

```

```

In [29]: # Categorical one-hot
cat_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore')) #we ignore the unknown cathe
])

# Numerical
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='median'))
])

# Ordinal
#age ordinal pipeline
age_pipeline = Pipeline([
    ('mapper', FunctionTransformer(
        lambda x: x.iloc[:, 0].map(age_mapping).to_frame(), #anonymous function, th
        validate=False)), #the mapping we have pr
    ('imputer', SimpleImputer(strategy='most_frequent'))
])

#total_units ordinal pipeline
total_units_pipeline = Pipeline([
    ('mapper', FunctionTransformer(
        lambda x: x.iloc[:, 0].map(total_units_mapping).to_frame(),
        validate=False)), #since we want to work with DataFrame we must set validat
    ('imputer', SimpleImputer(strategy='most_frequent'))
])

#Full pipeline, merging all type of data
preprocessor = ColumnTransformer([

```



```

random_state=42,
stratify=y, #in order to keep t
)

#stratified K-Fold for balanced class distribution over the target
strat_kfold = StratifiedKFold(n_splits=5, #number of folds
                              shuffle=True, #before creating the fold, it shuffles
                              random_state=42, #footnote following block
                              )

```

Notice: our dataset is a bit unbalanced in terms of target values, since the samples with **Status=0** are way more than the ones with **Status=1**, this can lead to slightly biased models. In order to overcome this issue we tweak some models' parameters accordingly.

This practice is especially useful for improving the **recall** of our models.

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

As opposed to:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Since credit institutes tend to be risk adverse in front of potential loan clients, they **prefer recall over precision**, and this is the way we approach our three following models.

random_state = 42 is an inside joke of machine learning based on the famous book: *The Hitchhiker's Guide to the Galaxy*, Douglas Adams, [see](#)

4.1. eXtreme Gradient Boosting

```

In [34]: #counting the number of positive and negative samples for balancing the model
num_negatives = loan_data['Status'].value_counts()[0]
num_positives = loan_data['Status'].value_counts()[1]

```

```

In [35]: """
Notice, if you do not have enough time, you can skip directly to the following cell
object in your folder, since it will open the already trained model.
"""

#setting the XGBoost classifier
xgb_model = XGBClassifier(eval_metric='logloss', #cost function
                          random_state=42, #to control the internal random componen
                          scale_pos_weight = num_negatives / num_positives #it give
                          )                                                    #since t
                                                                              #unbalan

#adding to the pipeline the algorithm
full_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', xgb_model)
])

```

```

])

#parameter grid for GridSearch
param_grid = {
    'classifier__n_estimators': [50, 100, 200],
    'classifier__max_depth': [5, 6, 7],
    'classifier__learning_rate': [0.1, 0.3, 0.5],
    'classifier__subsample': [0.8, 1], #fraction of samples used for fitting each
}

# Grid search setup
grid_search_XGBoost = GridSearchCV(full_pipeline,
                                    param_grid,
                                    cv=strat_kfold,
                                    n_jobs=-1, #it uses all the CPU cores availbale in paral
                                    verbose=2, #it shows all details for each combination
                                    scoring = 'recall', #since we want prioritize recall
                                    #scoring = 'f1',    #even this one Leads very good resul
                                    )

start_time = time.time() #for counting the time
# Fit GridSearchCV on training data
grid_search_XGBoost.fit(X_train,
                        y_train
                        )
end_time = time.time() #end time of execution
elapsed_time = convert_time(end_time - start_time)
print(f"XGBoosting training + grid search took {elapsed_time}")

# Best params and score
print("Best parameters found:", grid_search_XGBoost.best_params_)
print("Best cross-validation Recall:", grid_search_XGBoost.best_score_)

# Optional: Evaluate on the test set
test_score = grid_search_XGBoost.score(X_test, y_test)
print(f"Test set Recall: {test_score:.4f}") #we print the accuracy with the four de

```

Fitting 5 folds for each of 54 candidates, totalling 270 fits

XGBoosting training + grid search took 0:16:58

Best parameters found: {'classifier__learning_rate': 0.5, 'classifier__max_depth': 5, 'classifier__n_estimators': 50, 'classifier__subsample': 0.8}

Best cross-validation Recall: 0.7280060652009098

Test set Recall: 0.7481

```

In [36]: #try to see whether the above cell has been run
try:
    grid_search_XGBoost
    recovered = False #boolean variable to understand wheter the model has been rec

#otherwise open the serialized model
except:
    with open("XGBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary
        grid_search_XGBoost = cloudpickle.load(f)
    recovered = True

```

```

In [37]: pd.DataFrame(grid_search_XGBoost.cv_results_)

```

Out[37]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_learning
0	9.959908	0.111911	0.802918	0.367341	
1	12.227372	3.899827	1.755331	0.328310	
2	23.036407	2.366370	1.297397	0.365357	
3	14.486458	1.880278	0.800972	0.278468	
4	17.239578	0.220814	0.870704	0.023898	
5	13.535185	0.175856	0.833076	0.028101	
6	7.789968	0.249640	0.629102	0.022814	
7	6.889371	0.064463	0.632961	0.054496	
8	11.090194	0.225003	0.715223	0.050741	
9	9.592282	0.045793	0.728181	0.063756	
10	18.012373	0.280872	0.902797	0.023970	
11	14.426612	0.122053	0.922483	0.030842	
12	8.038273	0.083363	0.695766	0.046859	
13	8.249799	0.644209	0.918591	0.184402	
14	20.633695	0.349632	1.220312	0.245398	
15	19.012687	0.470595	1.817459	0.784448	
16	27.048638	2.356832	1.199594	0.087897	
17	16.867118	1.443681	1.051921	0.073208	
18	7.360573	0.173280	0.609933	0.043352	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_learning.
19	6.379111	0.076008	0.625277	0.061071	
20	10.775345	0.145443	0.721807	0.046877	
21	8.764243	0.168277	0.764832	0.094291	
22	16.549371	0.394369	0.845649	0.029795	
23	13.160515	0.536135	0.845918	0.028440	
24	7.339692	0.071558	0.615039	0.022919	
25	6.421416	0.075103	0.631852	0.025604	
26	10.895021	0.267455	0.732253	0.027924	
27	9.400118	0.199299	0.739972	0.036621	
28	19.909610	1.870109	0.994584	0.084985	
29	16.780405	0.524896	1.108583	0.126493	
30	8.864072	0.614887	0.712018	0.032819	
31	7.547914	0.119703	0.742236	0.049442	
32	16.395838	3.965111	1.662178	0.612789	
33	15.784762	3.779051	0.977600	0.050874	
34	23.934121	0.578148	1.415445	0.218974	
35	20.852833	1.539357	1.450688	0.298924	
36	9.749824	0.915292	0.951257	0.210750	
37	7.901964	0.537534	0.726036	0.071961	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_learning
38	12.846456	0.619731	0.952467	0.215300	
39	11.274114	1.802516	0.901988	0.160800	
40	27.085224	1.745152	1.155628	0.145150	
41	15.034133	1.750489	0.927211	0.134312	
42	7.424155	0.068395	0.653203	0.059335	
43	6.470957	0.096106	0.635905	0.037022	
44	10.870731	0.155201	0.740229	0.048882	
45	8.945145	0.131690	0.698178	0.015774	
46	18.597185	0.646536	0.917283	0.048571	
47	17.403107	2.578652	2.432832	1.788268	
48	8.154314	0.118954	0.729547	0.076788	
49	6.850456	0.116967	0.663318	0.039920	
50	11.778309	0.121604	0.783455	0.049557	
51	9.747708	0.058044	0.772599	0.062101	
52	20.359479	0.383227	1.114856	0.112094	
53	15.291070	1.323871	0.957885	0.215121	

```
In [38]: #we encapsulate just the best model
XGBoost_whole_model = grid_search_XGBoost.best_estimator_ #here we save the whole b
XGBoost_classifier = XGBoost_whole_model.named_steps['classifier'] #here we save ju
```

```
In [39]: #recover just the preprocessor from the whole model
preprocessor = XGBoost_whole_model.named_steps['preprocessor']
feature_names = []
```

```

for name, transformer, cols in preprocessor.transformers_:
    #if the transformer is a drop operator
    if transformer == 'drop':
        continue
    #if the transformer is actually a pipeline object, so that has inside other tra
    if isinstance(transformer, Pipeline):
        #get the last step of the pipeline, because the last step is the one that f
        last_step = transformer.steps[-1][1]
        if hasattr(last_step, 'get_feature_names_out'):
            names = last_step.get_feature_names_out(cols)
        else:
            names = cols
    #if the transformer is directly a transformer
    else:
        if hasattr(transformer, 'get_feature_names_out'):
            names = transformer.get_feature_names_out(cols)
        else:
            names = cols

    feature_names.extend(names)

feat_imp = pd.DataFrame({
    'feature': feature_names,
    'importance': XGBoost_classifier.feature_importances_
}).sort_values(
    by="importance",
    ascending=False
)

print(feat_imp)

```

	feature	importance
12	lump_sum_payment_lpsm	0.136860
43	property_value	0.125512
34	credit_type_CIB	0.121303
8	Neg_ammortization_neg_amm	0.079789
46	LTV	0.054352
18	submission_of_application_not_inst	0.053599
4	Credit_Worthiness_l1	0.034651
26	loan_type_type3	0.031678
47	dtir1	0.029537
24	loan_type_type1	0.027752
25	loan_type_type2	0.027425
2	approv_in_adv_nopre	0.021031
0	loan_limit_cf	0.020551
10	interest_only_int_only	0.017750
27	loan_purpose_p1	0.016836
44	income	0.016434
49	total_units	0.015406
32	occupancy_type_pr	0.015095
21	Gender_Joint	0.013179
28	loan_purpose_p2	0.013005
29	loan_purpose_p3	0.012184
37	Region_North	0.011764
14	construction_type_mh	0.010875
31	occupancy_type_ir	0.010173
41	loan_amount	0.010054
30	loan_purpose_p4	0.009058
16	co-applicant_credit_type_CIB	0.008699
33	occupancy_type_sr	0.005896
42	term	0.005726
48	age	0.005487
39	Region_central	0.005393
35	credit_type_CRIF	0.004872
23	Gender_Sex Not Available	0.004803
6	open_credit_nopc	0.004096
22	Gender_Male	0.004086
45	Credit_Score	0.003525
40	Region_south	0.003518
20	Gender_Female	0.003083
36	credit_type_EXP	0.002692
38	Region_North-East	0.002272
15	construction_type_sb	0.000000
9	Neg_ammortization_not_neg	0.000000
3	approv_in_adv_pre	0.000000
5	Credit_Worthiness_l2	0.000000
17	co-applicant_credit_type_EXP	0.000000
7	open_credit_opc	0.000000
19	submission_of_application_to_inst	0.000000
1	loan_limit_ncf	0.000000
11	interest_only_not_int	0.000000
13	lump_sum_payment_not_lpsm	0.000000

```
In [40]: # Predict on the test set
y_pred = grid_search_XGBoost.predict(X_test)

# Compute confusion matrix
```

```

cm = confusion_matrix(y_test, y_pred)

# Get predicted probabilities for the positive class
y_scores = grid_search_XGBoost.predict_proba(X_test)[: , 1]

# Compute average precision (AUC-PR)
auc_pr = average_precision_score(y_test, y_scores)

# Create a figure
fig, axes = plt.subplots(1, # one row
                        2, # two columns
                        figsize=(12, 4)
                        )

# --- Confusion Matrix ---
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=grid_search_XGBoost.classes_
                              )



disp.plot(cmap="Reds",
          values_format='d',
          colorbar=False, #don't show the legend colormap
          ax=axes[0])
axes[0].set_title("Confusion Matrix - XGBoost")

# --- Precision-Recall Curve ---
PrecisionRecallDisplay.from_predictions(y_test,
                                       y_scores,
                                       name="XGBoost",
                                       ax=axes[1]
                                       )

axes[1].set_title(f"Precision-Recall Curve - XGBoost (AUC={auc_pr:.4f})")

plt.tight_layout()
plt.show()

```



Notice: we kept the same balance of the original dataset based on target values for our test set, thanks to the parameter: **stratify=y** during the split of the dataset in section 4.. Thus, the preportion of: $\frac{\text{negatives}}{\text{positives}}=3.05$, is kept on both training and testing sets.


```
In [41]: # Predict on test set
y_pred = grid_search_XGBoost.predict(X_test)

# Accuracy
acc = accuracy_score(y_test, y_pred)
# Precision (by default, for positive class in binary classification)
prec = precision_score(y_test, y_pred)
# Recall
rec = recall_score(y_test, y_pred)
# F1-score
f1 = f1_score(y_test, y_pred)


print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1-score: {f1:.4f}")
```

Accuracy: 0.8692
Precision: 0.7286
Recall: 0.7481
F1-score: 0.7382

```
In [42]: # Get predicted probabilities for the positive class
y_scores = grid_search_XGBoost.predict_proba(X_test)[:, 1]

# Compute precision, recall, thresholds
precision, recall, thresholds = precision_recall_curve(y_test, y_scores)

# Plot Precision and Recall vs Threshold
plt.figure(figsize=(8, 3))
plt.plot(thresholds, precision[:-1], label='Precision', color='red')
plt.plot(thresholds, recall[:-1], label='Recall', color='blue')
plt.xlabel('Threshold')
plt.ylabel('Score')
plt.title('Precision & Recall vs Threshold - XGBoost')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [43]: y_scores = grid_search_XGBoost.predict_proba(X_test)[: , 1]
threshold = 0.4 # Lower than 0.5, where it is centered
y_pred_adjusted = (y_scores >= threshold).astype(int)
```


```
In [44]: # Compute confusion matrix
cm = confusion_matrix(y_test,
                      y_pred_adjusted
                      )

# Display confusion matrix as a heatmap
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=grid_search_XGBoost.classes_
                              )

plt.figure(figsize=(3, 3)) #create a specific figure object in order to better mani

disp.plot(cmap="Reds",
          values_format='d', #show numbers as integers
          colorbar = False, #colorbar as Legend disactivated
          ax=plt.gca() #plot inn the figure created
          )

plt.title("Confusion Matrix - XGBoost")
plt.show()
```



```
In [45]: # Accuracy
acc = accuracy_score(y_test, y_pred_adjusted)
# Precision (by default, for positive class in binary classification)
prec = precision_score(y_test, y_pred_adjusted)
# Recall
rec = recall_score(y_test, y_pred_adjusted)
# F1-score
f1 = f1_score(y_test, y_pred_adjusted)

print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
```

```
print(f"Recall: {rec:.4f}")
print(f"F1-score: {f1:.4f}")
```

Accuracy: 0.8255
Precision: 0.6117
Recall: 0.7997
F1-score: 0.6932

The following code block performs *serialization*, it is a very helpful practice that allows to save a model once has been trained, so it saves the training time every time you open the notebook by simply opening directly the saved model.

There are many libraries that apply serialization and probably one of the best-known is **pickle**, even though it is not able to save **lambda** functions when present.

Since this practice can also save the full pipeline, and in our case there are **lambda** functions, we ended up using **cloudpickle**, that also saves the anonymous functions. The storing format is binary.

```
In [46]: # --- Saving ---
if recovered is False: #if the model currently used has not been recovered from ser
    with open("XGBoost_loan.pkl", mode="wb") as f: #wb stands for Writing Binary
        cloudpickle.dump(grid_search_XGBoost, f)

# --- Loading ---
"""
with open("XGBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary
    XGBoost_recovered = cloudpickle.load(f)
"""
```

```
Out[46]: '\nwith open("XGBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary\n
XGBoost_recovered = cloudpickle.load(f)\n'
```

4.2. Adaptive Boosting

```
In [47]: """
Notice, if you do not have enough time, you can skip directly to the following cell
object in your folder, since it will open the already trained model.
"""

#Setting AdaBoost classifier (with Decision Tree as base estimator)
ada_model = AdaBoostClassifier(
    base_estimator=DecisionTreeClassifier(class_weight=
                                          ),
    random_state=42
)

#adding to the pipeline the algorithm
full_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', ada_model)
])

# Parameter grid for GridSearch
param_grid = {
```

```

    'classifier__base_estimator__max_depth': [3, 4, 5],          #Maximum depth of dec
    'classifier__n_estimators': [150, 200, 300],              #Number of boosting sta
    'classifier__learning_rate': [0.01, 0.3, 0.1],           #Shrinkage (learning) rate
    'classifier__algorithm': ['#SAMME', # Boosting algorithms, for our case SAMME.R
                              'SAMME.R']
}

# Grid search setup
grid_search_AdaBoost = GridSearchCV(
    full_pipeline,
    param_grid,
    cv=strat_kfold,
    n_jobs=-1,
    verbose=2,
    scoring='recall',
)

start_time = time.time() #for counting the time
# Fit GridSearchCV on training data
grid_search_AdaBoost.fit(X_train,
                        y_train
)
end_time = time.time() #end time of execution
elapsed_time = convert_time(end_time - start_time)
print(f"AdaBoost training + grid search took {elapsed_time}")

# Best params and score
print("Best parameters found:", grid_search_AdaBoost.best_params_)
print("Best cross-validation Recall:", grid_search_AdaBoost.best_score_)

# Evaluate on the test set
test_score = grid_search_AdaBoost.score(X_test, y_test)
print(f"Test set Recall: {test_score:.4f}")

```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

AdaBoost training + grid search took 4:45:10

Best parameters found: {'classifier__algorithm': 'SAMME.R', 'classifier__base_estimator__max_depth': 3, 'classifier__learning_rate': 0.3, 'classifier__n_estimators': 300}

Best cross-validation Recall: 0.7224867323730099

Test set Recall: 0.7424

```

In [48]: #try to see whether the above cell has been run
try:
    grid_search_AdaBoost
    recovered = False #boolean variable to understand wheter the model has been rec

#otherwise open the serialized model
except:
    with open("AdaBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary
        grid_search_AdaBoost = cloudpickle.load(f)
    recovered = True

```

```

In [49]: pd.DataFrame(grid_search_AdaBoost.cv_results_)

```

Out[49]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier__algorithm
0	245.091279	1.433680	4.111456	0.086499	SAMME.
1	317.415425	3.034288	5.322965	0.235682	SAMME.
2	472.035136	1.909110	7.570289	0.225191	SAMME.
3	238.743706	3.625973	4.291338	0.575546	SAMME.
4	364.822918	22.906856	6.342026	1.683179	SAMME.
5	466.482349	30.795482	6.480300	1.756797	SAMME.
6	244.422789	1.004910	4.046438	0.075939	SAMME.
7	321.772150	0.887950	5.266049	0.118813	SAMME.
8	515.447549	12.489158	10.032569	2.061255	SAMME.
9	332.836284	18.517381	3.979810	0.052671	SAMME.
10	410.594428	2.939652	5.124905	0.099056	SAMME.
11	653.750121	23.469345	7.898930	0.200846	SAMME.
12	307.835304	2.912320	3.975472	0.033814	SAMME.
13	446.468203	13.472946	5.554764	0.374650	SAMME.
14	634.295039	12.053299	7.712355	0.069699	SAMME.
15	322.399766	16.085456	4.129928	0.146897	SAMME.
16	438.795249	14.444886	6.547107	1.064233	SAMME.
17	638.956153	12.908956	8.337363	1.325155	SAMME.
18	406.464667	18.696185	4.304856	0.354997	SAMME.

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_algorithm
19	627.692060	33.060848	8.225628	0.734167	SAMME.
20	1103.511735	115.028874	13.411554	1.452328	SAMME.
21	654.125297	30.348780	7.687926	1.300510	SAMME.
22	624.216390	59.971174	5.803365	0.545394	SAMME.
23	720.234004	18.568695	7.333406	0.262206	SAMME.
24	345.514080	1.678799	3.768399	0.040220	SAMME.
25	463.378268	1.164982	4.973119	0.184906	SAMME.
26	617.579434	110.591147	6.869373	1.778494	SAMME.

```
In [50]: #we encapsulate just the best model
AdaBoost_whole_model = grid_search_AdaBoost.best_estimator_ #here we save the whole
AdaBoost_classifier = AdaBoost_whole_model.named_steps['classifier'] #here we save
```

```
In [51]: #recover just the preprocessor from the whole model
preprocessor = AdaBoost_whole_model.named_steps['preprocessor']
feature_names = []

for name, transformer, cols in preprocessor.transformers_:
    #if the transformer is a drop operator
    if transformer == 'drop':
        continue
    #if the transformer is actually a pipeline object, so that has inside other tra
    if isinstance(transformer, Pipeline):
        #get the last step of the pipeline, because the last step is the one that f
        last_step = transformer.steps[-1][1]
        if hasattr(last_step, 'get_feature_names_out'):
            names = last_step.get_feature_names_out(cols)
        else:
            names = cols
    #if the transformer is directly a transformer
    else:
        if hasattr(transformer, 'get_feature_names_out'):
            names = transformer.get_feature_names_out(cols)
        else:
            names = cols

    feature_names.extend(names)

feat_imp = pd.DataFrame({
```

```
        'feature': feature_names,  
        'importance': AdaBoost_classifier.feature_importances_  
    }).sort_values(  
        by="importance",  
        ascending=False  
    )  
  
print(feat_imp)
```

	feature	importance
46	LTV	0.205134
44	income	0.131380
43	property_value	0.098443
47	dtir1	0.090626
41	loan_amount	0.078529
45	Credit_Score	0.065378
29	loan_purpose_p3	0.024724
48	age	0.024074
31	occupancy_type_ir	0.022039
42	term	0.021247
27	loan_purpose_p1	0.018398
24	loan_type_type1	0.017426
30	loan_purpose_p4	0.017361
26	loan_type_type3	0.011771
2	approv_in_adv_nopre	0.010662
22	Gender_Male	0.009872
21	Gender_Joint	0.008940
9	Neg_ammortization_not_neg	0.008019
49	total_units	0.008008
25	loan_type_type2	0.007877
19	submission_of_application_to_inst	0.007611
8	Neg_ammortization_neg_amm	0.007359
17	co-applicant_credit_type_EXP	0.007241
4	Credit_Worthiness_l1	0.007118
18	submission_of_application_not_inst	0.006882
3	approv_in_adv_pre	0.006844
1	loan_limit_ncf	0.005524
37	Region_North	0.005429
13	lump_sum_payment_not_lpsm	0.005085
36	credit_type_EXP	0.005038
5	Credit_Worthiness_l2	0.004803
39	Region_central	0.004575
28	loan_purpose_p2	0.004568
16	co-applicant_credit_type_CIB	0.004443
10	interest_only_int_only	0.004397
34	credit_type_CIB	0.004236
35	credit_type_CRIF	0.004229
32	occupancy_type_pr	0.004019
12	lump_sum_payment_lpsm	0.003556
33	occupancy_type_sr	0.003234
20	Gender_Female	0.003051
11	interest_only_not_int	0.002444
0	loan_limit_cf	0.002260
23	Gender_Sex Not Available	0.001560
6	open_credit_nopc	0.001486
7	open_credit_opc	0.001085
40	Region_south	0.001079
15	construction_type_sb	0.000562
38	Region_North-East	0.000374
14	construction_type_mh	0.000000

```
In [52]: # Predict on the test set
y_pred = grid_search_AdaBoost.predict(X_test)

# Compute confusion matrix
```



```

cm = confusion_matrix(y_test, y_pred)

# Get predicted probabilities for the positive class
y_scores = grid_search_AdaBoost.predict_proba(X_test)[: , 1]

# Compute average precision (AUC-PR)
auc_pr = average_precision_score(y_test, y_scores)

# Create a figure
fig, axes = plt.subplots(1, # one row
                        2, # two columns
                        figsize=(12, 4)
                        )

# --- Confusion Matrix ---
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=grid_search_AdaBoost.classes_
                              )



disp.plot(cmap="Reds",
          values_format='d',
          colorbar=False, #don't show the legend colormap
          ax=axes[0])
axes[0].set_title("Confusion Matrix - AdaBoost")

# --- Precision-Recall Curve ---
PrecisionRecallDisplay.from_predictions(y_test,
                                       y_scores,
                                       name="AdaBoost",
                                       ax=axes[1]
                                       )

axes[1].set_title(f"Precision-Recall Curve - AdaBoost (AUC={auc_pr:.4f})")

plt.tight_layout()
plt.show()

```



```

In [53]: # Predict on test set
y_pred = grid_search_AdaBoost.predict(X_test)

# Accuracy
acc = accuracy_score(y_test, y_pred)
# Precision (by default, for positive class in binary classification)
prec = precision_score(y_test, y_pred)

```

```

# Recall
rec = recall_score(y_test, y_pred)
# F1-score
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1-score: {f1:.4f}")

```

Accuracy: 0.8721
Precision: 0.7395
Recall: 0.7424
F1-score: 0.7409


```

In [54]: # Get predicted probabilities for the positive class
y_scores = grid_search_AdaBoost.predict_proba(X_test)[:, 1]

# Compute precision, recall, thresholds
precision, recall, thresholds = precision_recall_curve(y_test, y_scores)

# Plot Precision and Recall vs Threshold
plt.figure(figsize=(8, 3))
plt.plot(thresholds, precision[:-1], label='Precision', color='red')
plt.plot(thresholds, recall[:-1], label='Recall', color='blue')
plt.xlabel('Threshold')
plt.ylabel('Score')
plt.title('Precision & Recall vs Threshold - AdaBoost')
plt.legend()
plt.grid(True)
plt.show()

```



```

In [110... y_scores = grid_search_AdaBoost.predict_proba(X_test)[:, 1]
threshold = 0.499 # Lower than 0.5, where it is centered
y_pred_adjusted = (y_scores >= threshold).astype(int)

```

```

In [111... # Compute confusion matrix
cm = confusion_matrix(y_test,
                      y_pred_adjusted)

```

```

    )


# Display confusion matrix as a heatmap
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=grid_search_AdaBoost.classes_)

plt.figure(figsize=(3, 3)) #create a specific figure object in order to better mani

disp.plot(cmap="Reds",
          values_format='d', #show numbers as integers
          colorbar = False, #colorbar as legend disactivated
          ax=plt.gca() #plot inn the figure created
        )

plt.title("Confusion Matrix - AdaBoost")
plt.show()

```



In [112...

```

# Accuracy
acc = accuracy_score(y_test, y_pred_adjusted)
# Precision (by default, for positive class in binary classification)
prec = precision_score(y_test, y_pred_adjusted)
# Recall
rec = recall_score(y_test, y_pred_adjusted)
# F1-score
f1 = f1_score(y_test, y_pred_adjusted)

print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1-score: {f1:.4f}")

```

Accuracy: 0.8326
Precision: 0.6280
Recall: 0.7868
F1-score: 0.6985

```
In [58]: # --- Saving ---
if recovered is False: #if the model currently used has not been recovered from ser
    with open("AdaBoost_loan.pkl", mode="wb") as f: #wb stands for Writing Binary
        cloudpickle.dump(grid_search_AdaBoost, f)

# --- Loading ---
"""
with open("AdaBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary
    AdaBoost_recovered = cloudpickle.load(f)
"""

Out[58]: '\nwith open("AdaBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary\n
AdaBoost_recovered = cloudpickle.load(f)\n'
```

4.3. Gradient Boosting

```
In [59]: """
Notice, if you do not have enough time, you can skip directly to the following cell
object in your folder, since it will open the already trained model.
"""

#setting the GradientBoosting model
gb_model = GradientBoostingClassifier(
    random_state=42
)

#adding to the pipeline the algorithm
full_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', gb_model)
])

# Parameter grid for GridSearch
param_grid = {
    'classifier__n_estimators': [250, 350, 400], #Number of boosting stage
    'classifier__learning_rate': [0.1, 0.2], #Shrinkage (learning) rate
    'classifier__max_depth': [5, 6], #Maximum depth of decision tr
    'classifier__subsample': [0.8, 1.0], #fraction of samples used
    'classifier__min_samples_split': [7, 10] #minimum samples required
}

# Grid search setup
grid_search_GradientBoost = GridSearchCV(
    full_pipeline,
    param_grid,
    cv=StratifiedKFold,
    n_jobs=-1,
    verbose=2,
    scoring='recall',
)

start_time = time.time() #for counting the time
# Fit GridSearchCV on training data
grid_search_GradientBoost.fit(X_train,
```

```

        y_train
    )
end_time = time.time() #end time of execution
elapsed_time = convert_time(end_time - start_time)
print(f"GradientBoost training + grid search took {elapsed_time}")

# Best params and score
print("Best parameters found:", grid_search_GradientBoost.best_params_)
print("Best cross-validation Recall:", grid_search_GradientBoost.best_score_)

# Optional: Evaluate on the test set
test_score = grid_search_GradientBoost.score(X_test, y_test)
print(f"Test set Recall: {test_score:.4f}")

```

Fitting 5 folds for each of 48 candidates, totalling 240 fits
 GradientBoost training + grid search took 11:53:26
 Best parameters found: {'classifier__learning_rate': 0.2, 'classifier__max_depth': 6, 'classifier__min_samples_split': 7, 'classifier__n_estimators': 400, 'classifier__subsample': 0.8}
 Best cross-validation Recall: 0.6412736921910538
 Test set Recall: 0.6638

```

In [60]: #try to see whether the above cell has been run
try:
    grid_search_GradientBoost
    recovered = False #boolean variable to understand wheter the model has been rec

#otherwise open the serialized model
except:
    with open("GradientBoost_loan.pkl", mode="rb") as f: #rb stands for reading bin
        grid_search_GradientBoost = cloudpickle.load(f)
    recovered = True

```

```

In [61]: pd.DataFrame(grid_search_GradientBoost.cv_results_)

```

Out[61]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_learning
0	512.879887	144.450418	1.122826	0.063351	
1	529.986002	5.964861	0.987681	0.024942	
2	600.106107	3.549016	1.210050	0.031430	
3	739.277332	3.111717	1.232336	0.033994	
4	685.497165	2.432012	1.357316	0.045910	
5	847.365082	2.331266	1.351080	0.040197	
6	430.180005	2.764529	0.984268	0.015888	
7	517.996707	9.113573	1.025322	0.039395	
8	598.396039	1.338019	1.217375	0.010686	
9	735.358536	1.416284	1.214518	0.016678	
10	683.282903	0.908090	1.346113	0.030429	
11	844.259007	3.414119	1.331821	0.023504	
12	507.501319	0.807878	1.027414	0.021208	
13	616.086510	9.216419	1.039808	0.039840	
14	710.775291	1.752213	1.285066	0.008279	
15	884.544131	5.831086	1.295912	0.018324	
16	819.894631	7.279384	1.446876	0.050401	
17	1010.058585	7.963290	1.413886	0.011855	
18	509.259252	4.560607	1.040331	0.016265	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_learning.
19	627.550477	3.987945	1.047841	0.016237	
20	711.698959	2.387991	1.296814	0.008935	
21	884.885003	2.905104	1.329976	0.066586	
22	821.041016	3.795591	1.429279	0.022201	
23	985.777877	14.838792	1.433254	0.032422	
24	434.285538	4.884939	1.011667	0.011140	
25	537.809077	1.623921	1.045920	0.017239	
26	611.762036	1.573480	1.265683	0.015563	
27	752.261048	1.697127	1.257984	0.009191	
28	698.959441	1.681796	1.611785	0.471492	
29	841.966967	9.355861	1.373461	0.010924	
30	438.581219	2.159170	1.023608	0.005995	
31	538.258345	0.487609	1.015282	0.015603	
32	613.360212	2.590131	1.257805	0.024218	
33	754.772213	3.014548	1.265249	0.016140	
34	694.920612	2.101590	1.362522	0.015444	
35	835.277452	13.329384	1.383765	0.037637	
36	518.270203	3.715741	1.067674	0.014647	
37	634.765384	1.988271	1.068276	0.022410	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_classifier_learning_
38	714.320291	3.112515	1.312731	0.013647	
39	890.356776	3.208117	1.313919	0.030577	
40	818.291143	2.011892	1.456371	0.020871	
41	1017.280905	2.839392	1.437010	0.032551	
42	509.734762	1.242842	1.065272	0.015862	
43	632.832738	2.010486	1.059929	0.023055	
44	715.706904	3.373535	1.347079	0.043799	
45	866.709132	14.490509	1.293606	0.008553	
46	811.790967	4.671743	1.456968	0.036732	
47	918.061253	124.464338	1.118934	0.291508	

```
In [62]: #we encapsulate just the best model
GradientBoost_whole_model = grid_search_GradientBoost.best_estimator_ #here we save
GradientBoost_classifier = GradientBoost_whole_model.named_steps['classifier'] #her
```

```
In [63]: #recover just the preprocessor from the whole model
preprocessor = GradientBoost_whole_model.named_steps['preprocessor']
feature_names = []

for name, transformer, cols in preprocessor.transformers_:
    #if the transformer is a drop operator
    if transformer == 'drop':
        continue
    #if the transformer is actually a pipeline object, so that has inside other tra
    if isinstance(transformer, Pipeline):
        #get the last step of the pipeline, because the last step is the one that f
        last_step = transformer.steps[-1][1]
        if hasattr(last_step, 'get_feature_names_out'):
            names = last_step.get_feature_names_out(cols)
        else:
            names = cols
    #if the transformer is directly a transformer
    else:
        if hasattr(transformer, 'get_feature_names_out'):
            names = transformer.get_feature_names_out(cols)
        else:
```



```
names = cols

feature_names.extend(names)

feat_imp = pd.DataFrame({
    'feature': feature_names,
    'importance': GradientBoost_classifier.feature_importances_
}).sort_values(
    by="importance",
    ascending=False
)

print(feat_imp)
```

	feature	importance
43	property_value	0.306946
46	LTV	0.254441
47	dtir1	0.080428
44	income	0.063137
45	Credit_Score	0.031656
34	credit_type_CIB	0.028567
41	loan_amount	0.026394
24	loan_type_type1	0.023346
12	lump_sum_payment_lpsm	0.015999
8	Neg_ammortization_neg_amm	0.013691
26	loan_type_type3	0.011565
48	age	0.009984
13	lump_sum_payment_not_lpsm	0.009810
27	loan_purpose_p1	0.009537
9	Neg_ammortization_not_neg	0.009407
4	Credit_Worthiness_l1	0.009114
19	submission_of_application_to_inst	0.007774
42	term	0.006949
3	approv_in_adv_pre	0.006822
25	loan_type_type2	0.006762
29	loan_purpose_p3	0.006734
18	submission_of_application_not_inst	0.005684
31	occupancy_type_ir	0.005153
30	loan_purpose_p4	0.004846
5	Credit_Worthiness_l2	0.003642
0	loan_limit_cf	0.003250
1	loan_limit_ncf	0.002686
2	approv_in_adv_nopre	0.002615
49	total_units	0.002610
32	occupancy_type_pr	0.002479
39	Region_central	0.002139
37	Region_North	0.002114
28	loan_purpose_p2	0.002072
21	Gender_Joint	0.002005
22	Gender_Male	0.001984
16	co-applicant_credit_type_CIB	0.001968
23	Gender_Sex Not Available	0.001961
35	credit_type_CRIF	0.001826
10	interest_only_int_only	0.001752
36	credit_type_EXP	0.001739
17	co-applicant_credit_type_EXP	0.001429
20	Gender_Female	0.001421
11	interest_only_not_int	0.001183
33	occupancy_type_sr	0.001170
40	Region_south	0.001095
38	Region_North-East	0.000806
14	construction_type_mh	0.000474
15	construction_type_sb	0.000463
7	open_credit_opc	0.000189
6	open_credit_nopc	0.000183

```
In [64]: # Predict on the test set
y_pred = grid_search_GradientBoost.predict(X_test)

# Compute confusion matrix
```

```

cm = confusion_matrix(y_test, y_pred)

# Get predicted probabilities for the positive class
y_scores = grid_search_GradientBoost.predict_proba(X_test)[:, 1]

# Compute average precision (AUC-PR)
auc_pr = average_precision_score(y_test, y_scores)

# Create a figure
fig, axes = plt.subplots(1, # one row
                        2, # two columns
                        figsize=(12, 4)
                        )

# --- Confusion Matrix ---
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=grid_search_GradientBoost.classes_
                              )



disp.plot(cmap="Reds",
          values_format='d',
          colorbar=False, #don't show the legend colormap
          ax=axes[0])
axes[0].set_title("Confusion Matrix - GradientBoost")

# --- Precision-Recall Curve ---
PrecisionRecallDisplay.from_predictions(y_test,
                                       y_scores,
                                       name="GradientBoost",
                                       ax=axes[1]
                                       )

axes[1].set_title(f"Precision-Recall Curve - GradientBoost (AUC={auc_pr:.4f})")

plt.tight_layout()
plt.show()

```



```

In [65]: # Predict on test set
y_pred = grid_search_GradientBoost.predict(X_test)

# Accuracy
acc = accuracy_score(y_test, y_pred)
# Precision (by default, for positive class in binary classification)
prec = precision_score(y_test, y_pred)

```

```

# Recall
rec = recall_score(y_test, y_pred)
# F1-score
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1-score: {f1:.4f}")

```

Accuracy: 0.8959
Precision: 0.8853
Recall: 0.6638
F1-score: 0.7587


```

In [66]: # Get predicted probabilities for the positive class
y_scores = grid_search_GradientBoost.predict_proba(X_test)[: , 1]

# Compute precision, recall, thresholds
precision, recall, thresholds = precision_recall_curve(y_test, y_scores)

# Plot Precision and Recall vs Threshold
plt.figure(figsize=(8, 3))
plt.plot(thresholds, precision[:-1], label='Precision', color='red')
plt.plot(thresholds, recall[:-1], label='Recall', color='blue')
plt.xlabel('Threshold')
plt.ylabel('Score')
plt.title('Precision & Recall vs Threshold - GradientBoost')
plt.legend()
plt.grid(True)
plt.show()

```



```

In [123... y_scores = grid_search_GradientBoost.predict_proba(X_test)[: , 1]
threshold = 0.21 # Lower than 0.5, where it is centered
y_pred_adjusted = (y_scores >= threshold).astype(int)

```

```

In [124... # Compute confusion matrix
cm = confusion_matrix(y_test,
                      y_pred_adjusted)

```

```

    )


# Display confusion matrix as a heatmap
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=grid_search_GradientBoost.classes_)

plt.figure(figsize=(3, 3)) #create a specific figure object in order to better mani

disp.plot(cmap="Reds",
          values_format='d', #show numbers as integers
          colorbar = False, #colorbar as legend disactivated
          ax=plt.gca() #plot inn the figure created
        )

plt.title("Confusion Matrix - GradientBoost")
plt.show()

```



In [125...

```

# Accuracy
acc = accuracy_score(y_test, y_pred_adjusted)
# Precision (by default, for positive class in binary classification)
prec = precision_score(y_test, y_pred_adjusted)
# Recall
rec = recall_score(y_test, y_pred_adjusted)
# F1-score
f1 = f1_score(y_test, y_pred_adjusted)

print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1-score: {f1:.4f}")

```

Accuracy: 0.8491
Precision: 0.6701
Recall: 0.7639
F1-score: 0.7139

```
In [70]: # --- Saving ---
if recovered is False: #if the model currently used has not been recovered from ser
    with open("GradientBoost_loan.pkl", mode="wb") as f: #wb stands for Writing Bin
        cloudpickle.dump(grid_search_GradientBoost, f)


# --- Loading ---
"""
with open("GradientBoost_loan.pkl", mode="rb") as f: #rb stands for reading binary
    GradientBoost_recovered = cloudpickle.load(f)
"""

Out[70]: '\nwith open("GradientBoost_loan.pkl", mode="rb") as f: #rb stands for reading bin
ary\n    GradientBoost_recovered = cloudpickle.load(f)\n'
```

5. Comparisons

Now that all three models have been run we can make comparisons and draw conclusions.

Since in our dataset the positive class is rare, as a rule of thumb we should prefer the *Precision-Recall* curves for comparisons, hence we will combine them with the best scores found after tweaking the precision-recall threshold, in order to choose the best model.



See [Appendix 7.B](#) to understand how these images have been generated

All three models are quite good, and have similar scores.

On the one hand XGBoost and GradientBoost have different scores each one with its own pros and cons, while AdaBoost seems to be the model in the middle.

If we take into account the overall F1-score that is basically the harmonic mean between precision and recall:

$$F1\text{-score} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

We would choose GradientBoost, since this metric is comprehensive of both Precision and Recall.

If we want to dig deeper into the curves, the shape of the PR-curve for GradientBoosting has a strange short drop at the very beginning, when the threshold is set at nearly its maximum. This phenomena could derive from several reasons, but mathematically speaking it depends on the fact that when we set the threshold at very high values, the number of samples captured is very low, so the **precision** can encompass only few true positive samples, so if the model misclassifies even a very low amount of negative instances, the metric goes down all of a sudden.

This is supported by the fact that our dataset, as previously shown, has a low amount of positive instances compared to negative ones.

Furthermore, the Precision & Recall Vs Threshold, curve for GradientBoosting and XGBoosting, demonstrate this artefact, because when the threshold is nearly at its maximum, in these two cases, the Recall drops all of a sudden.

Another possible explanation for the unusual behaviour of GradientBoost's PR-curve is that unlike the other two models, GradientBoost does not have a straightforward parameter to compensate an unbalanced dataset; we can do it, but in a longer and a more manual setup (*e.g.*: via sample weights). (Notice that for XGboost and AdaBoost, we set the models at the beginning with balancing out the dataset).

In any case, this is a common behaviour when dealing with unbalanced datasets, and does not necessarily mean that a model is weak, instead GradientBoost is again quite good, as the Area Under the Curve witnesses.

On the other hand if we take into account the general principle of this project, we want to prioritize recall over precision and even further, if we take a look at training time:

XGBoost	AdaBoost	GradientBoost
00:16:58	04:45:10	11:53:26

XGBoost is the one that performs better

Even under the perspective of the Area Under the Curve, that for comparisons is one of the main metrics that is taken under consideration, as it gives a numerical summary of a model's overall performance. Particularly useful for imbalanced datasets. It represents the total area enclosed by the PR-curve, which plots precision against recall across different classification thresholds. A higher PR AUC value, closer to 1, indicates a better-performing model, showing a superior ability to correctly identify positive cases while minimizing false positives.

All this is due to some intrinsic characteristics of the implementation of the algorithm itself. XGBoost is built upon the core Gradient Boosting framework, but it has some key features such as smarter penalizations and regularization, adopts Newton Boosting, furthermore XGBoost can build parts of the model in parallel. This not only speeds up the training process but also makes it scalable, allowing it to handle very large datasets by distributing

the workload across multiple machines or processing units.

Therefore, although GradientBoost achieves the best F1 and Accuracy, XGBoost provides the best trade-off between Recall, overall PR AUC, and computational efficiency, making it the most suitable choice for our project.

Notice: if your model takes a lot of time to train (e.g.: 24h+), make sure to periodically save the jupyter notebook file, because the Cross-site request forgery token (`XSRF`) may get out of sync, leading to a restart of the page without saving the variables! [see](#)

6. Conclusion and Possible Expansions

A simple and very effective strategy that we can implement in order to enhance quickly our models can be achieved by comparing the best parameters found for each model with the ranges that we have set in the `param_grid` dictionaries.

By doing so, we can tweak the ranges accordingly and re-run the models many times until we reach the optimal solutions with the proper ranges.

With a similar approach we can even expand the *hyperparameters* to tweak, since there exist many more, of course this will lead to slow down the process. A possible way to cope with this latter issue is by using `RandomizedSearchCV`, it evaluates a given number of random combinations by selecting a random value for each hyperparameter at every iteration.

This project can be further improved under some aspects that can be seen as a good starting point for future directions.

One of the first improvement that can come up in mind straight away after analyzing section 4. is the further feature selection, since all three algorithms under the hood of features importance show that: `open_credit_opc`, `open_credit_nopc` and `construction_type_sb` have no importance in reducing impurity.

Plotting the **learning curves** and eventually implementing some strategies such as **early stopping** can be very beneficial especially for Adaptive Boosting and Gradient Boosting, since they take long time for training.

Since the dataset results a bit unbalanced towards the negative samples, we could implement **SMOTE** (*Synthetic Minority Over-sampling Technique*) techniques in order to generate synthetic samples from the minority class, so to compensate it.

Furthermore we can try to apply dimensionality reduction and/or different models, maybe even approaching the field of Artificial Neural Networks, starting with the base model of the Multilayer Perceptron.

The main purpose of this notebook has originally started as a Kaggle challenge, it can be actually seen a good starting point for many other interesting applications, not necessarily

related to the credit industry.

For example with the same binary classifier concept we can implement fraud detectors, spam email detector, malicious cyber attacks identifier and so on...

7. Appendix

This section contains some tips and interesting tricks that I have used during the execution of this Notebook.

7.A. Python stay-awake module

At the end of section 5. we highlighted the training time for the three models. If we sum all training times we quickly understand that in total, running this notebook from scratch requires approximately seventeen hours.

If we do not want to change the settings of our computer in order to prevent sleep mode, we can use a Python script that runs in a different thread as a background daemon, since it requires very few CPU.


Online there are many modules from which we can take advantage of.

For example this one [here](#) is very easy and straightforward to implement, even though if you run it on Windows 10 or newer versions you need to adjust it a bit, because automatic mouse moving does not prevent sleep mode anymore [see](#).

Thus, once you have found out the location of the script in your computer with: `pip show stay-awake` you can change the `__init__.py` inside the `stay-awake` folder script so that instead of moving the mouse it activates/deactivates the ScrollLock button.

Once the script is ready, it is very easy to use it, just open a Shell prompt and type: `python -`

`m stay-awake`



```
Prompt dei comandi
Microsoft Windows [Versione 10.0.19045.6216]
(c) Microsoft Corporation. Tutti i diritti sono riservati.

C:\Users\Francesco>
C:\Users\Francesco>
C:\Users\Francesco>python -m stay-awake
>> Starting

>> -*- Thank You for using Stay-Awake, please STAR this repo at https://github.com/singhsidhukuldeep/stay-awake ! -*- <<

>> Importing Libraries
>> Libraries Imported
>> Setting Up
>> Wait time set to: 1 minutes
>> Setup Complete
>> Triggering stay-awake via Scroll Lock toggle
>> Scroll Lock toggled at 14:06:03.873251
>> Waiting for 60 seconds: 100% [ 60/60 [01:00<00:00, 1.00s/it]
>> Triggering stay-awake via Scroll Lock toggle
>> Scroll Lock toggled at 14:07:04.069019
>> Waiting for 60 seconds: 20% [ 12/60 [00:12<00:48, 1.00s/it]
C
C:\Users\Francesco>
```

See [Appendix 7.B.](#) to understand how this image has been generated

7.B. Python image converter base64 (PDF exporting)

If we want to convert a file Jupyter Notebook with images into PDF format, all the media attachments will be lost, this is a well-known bug of `ipynb` files.

In order to overcome this issue there are a few tricks, the one I have implemented is based on the Python library `base64`. Basically it converts images into ASCII characters, by opening the attachment in binary mode and then encoding it in series of printable characters limited to a set of 64 unique characters. More specifically, the source binary data is taken 6 bits at a time, then this group of 6 bits is mapped to one of 64 unique characters.

Base64 encoding causes an overhead of 33–37% relative to the size of the original binary data (33% by the encoding itself; up to 4% more by the inserted line breaks).

For further explanations see: [GitHub issue](#), [Base64](#)

The following block of codes are the ones that generated the encoding and stored directly the content into `HTML` tags for the images we saw through this notebook earlier.

Of course, once the attachment has been encoded, we can remove it from the directory.

For the sake of brevity we do not show the output of the blocks.

Each block is led by the `Markdown` referenced cell, that we would have written in case we had not had any issues.

```
In [ ]: 
```

```
In [137... import base64 #Library for base64 encoding

file = "loan.JPG"

with open(file, "rb") as f:
    data = f.read()

encoded = base64.b64encode(data).decode("utf-8")

# Genera il tag HTML con base64
html_code = f''

#print(html_code)
```

```
In [ ]: <div style="display: flex;
        justify-content: center;
        text-align: center;">

    <div>
        XGBoost
        <br>
        Accuracy: 0.8300 <br>
        Precision: 0.6105 <br>
        Recall: 0.7780 <br>
        F1-score: 0.6842
    </div>
```

```

<div>
    AdaBoost
    <br>
    Accuracy: 0.8247 <br>
    Precision: 0.6020 <br>
    Recall: 0.7652 <br>
    F1-score: 0.6738
</div>

<div>
    GradientBoost
    <br>
    Accuracy: 0.8444 <br>
    Precision: 0.6470 <br>
    Recall: 0.7533 <br>
    F1-score: 0.6961
</div>

</div>

```

In [134...

```

files = ["PR-curve XGBoost.png",
         "PR-curve AdaBoost.png",
         "PR-curve GradientBoost.png",
        ]
labels = ["XGBoost",
          "AdaBoost",
          "GradientBoost",
        ]
accuracy = [0.8255,
            0.8326,
            0.8491,
            ]
precision = [0.6117,
             0.6280,
             0.6701,
             ]
recall = [0.7997,
          0.7868,
          0.7639,
          ]
f1 = [0.6932,
      0.6985,
      0.7139,
      ]

html = '<div style="display: flex; justify-content: center; text-align: center;">\n

for i in range(len(files)):
    with open(files[i], "rb") as f:
        data = f.read()
        encoded = base64.b64encode(data).decode("utf-8")

    html += f''' <div>
        {labels[i]}
        <br>

```

```

        Accuracy: {accuracy[i]:.4f} <br>
        Precision: {precision[i]:.4f} <br>
        Recall: {recall[i]:.4f} <br>
        F1-score: {f1[i]:.4f}
    </div>\n\n'''

html += '</div>'

#print(html)

```

```
In [ ]: 
```

```
In [73]: file = "stay-awake.JPG"

with open(file, "rb") as f:
    data = f.read()

encoded = base64.b64encode(data).decode("utf-8")

html_code = f''

# print(html_code)

```